



LoRA Training Comparison: Forward vs Critical vs Creative

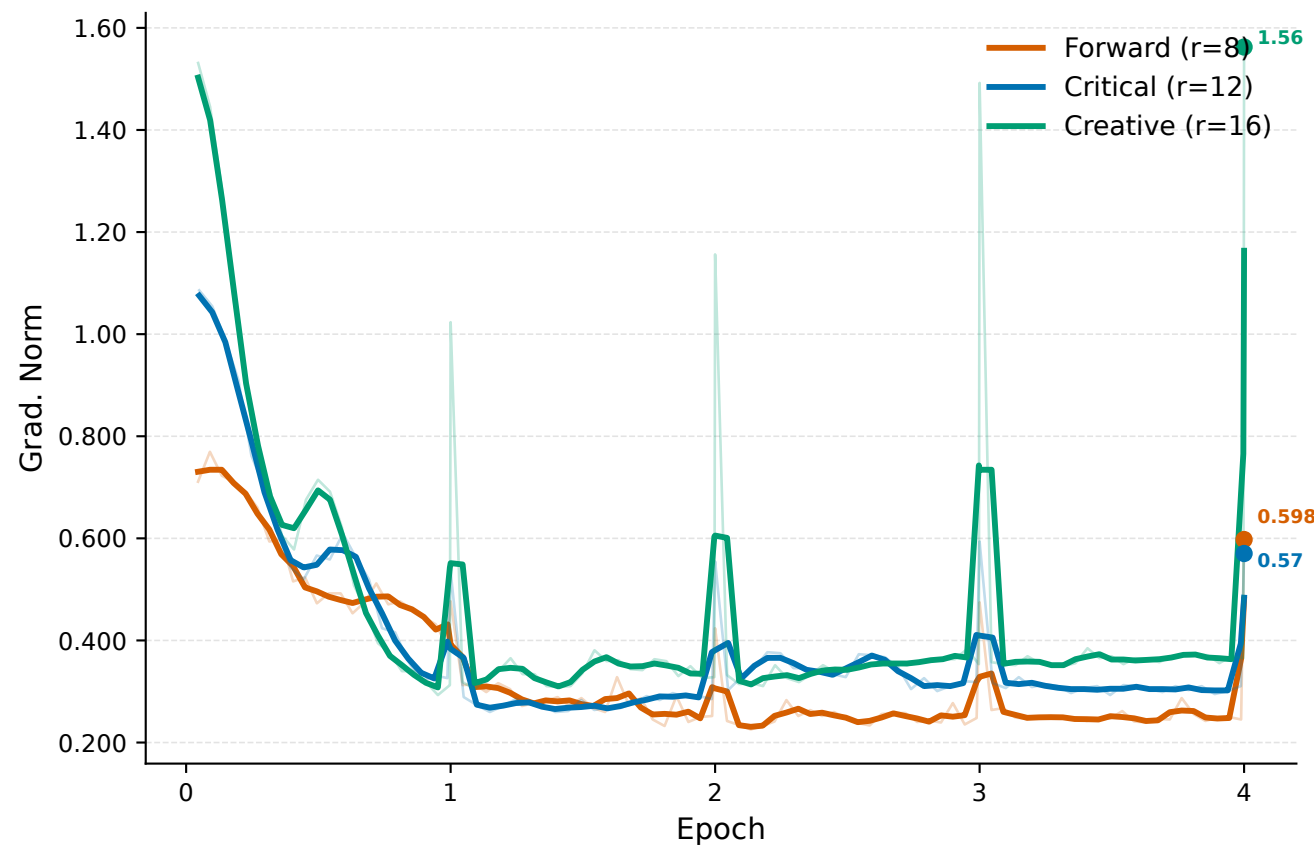
Training Loss



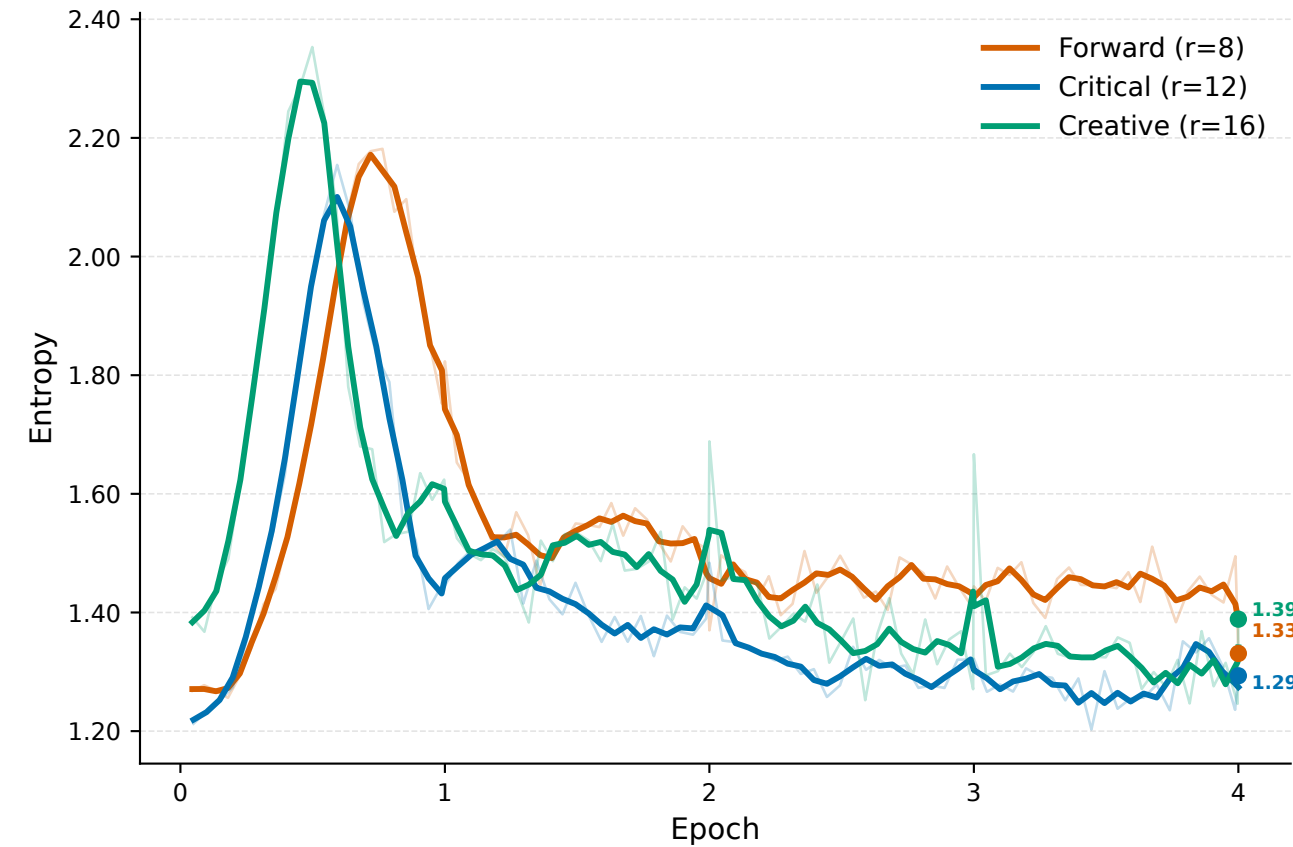
Mean Token Accuracy



Gradient Norm



Policy Entropy



Thin lines show raw values; bold lines show a 3-epoch moving average. Each model uses different LoRA rank and target modules.