



Critical LoRA Training Metrics


Training Dynamics Overview




Training Loss [1.177, 2.416]




Gradient Norm [0.2226, 2.168]



Policy Entropy [0.8335, 2.255]



Mean Token Accuracy [0.5186, 0.6622]



Learning Rate [0, 0.0002]

